

SWE3004 Operating Systems, fall 2023

Project 3. CPU Scheduling

TA)

Jongseok Kim

Shinhyun Park

Hyeonmyeong Lee

Gwanjong Park

Project plan

- Total 6 projects

- ~~1) Booting xv6 operating system~~

- ~~2) System call~~

- 3) CPU scheduling

- Linux CFS scheduler

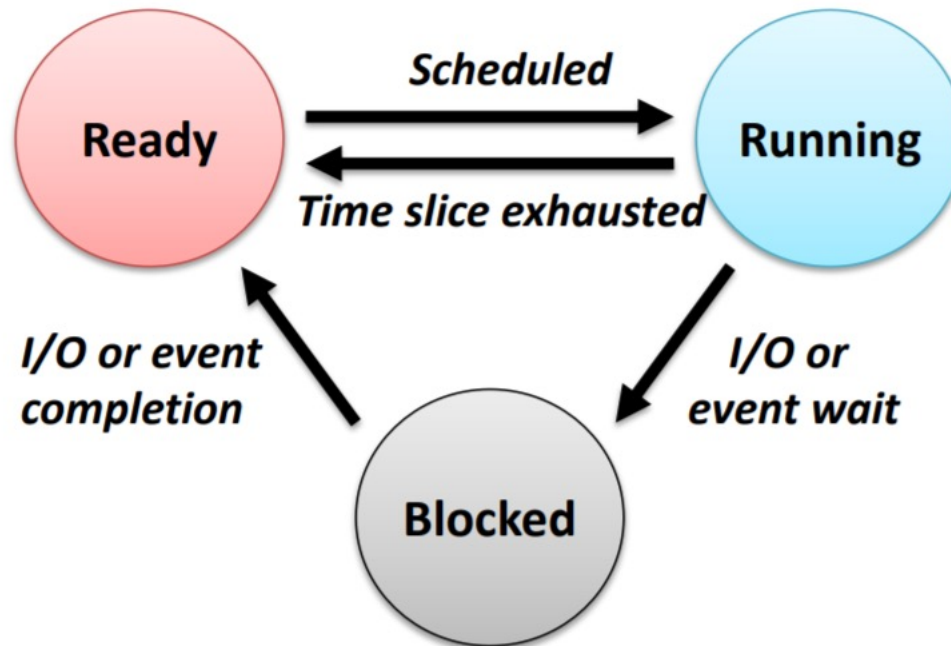
- 4) Virtual memory

- 5) Page replacement

- 6) File systems

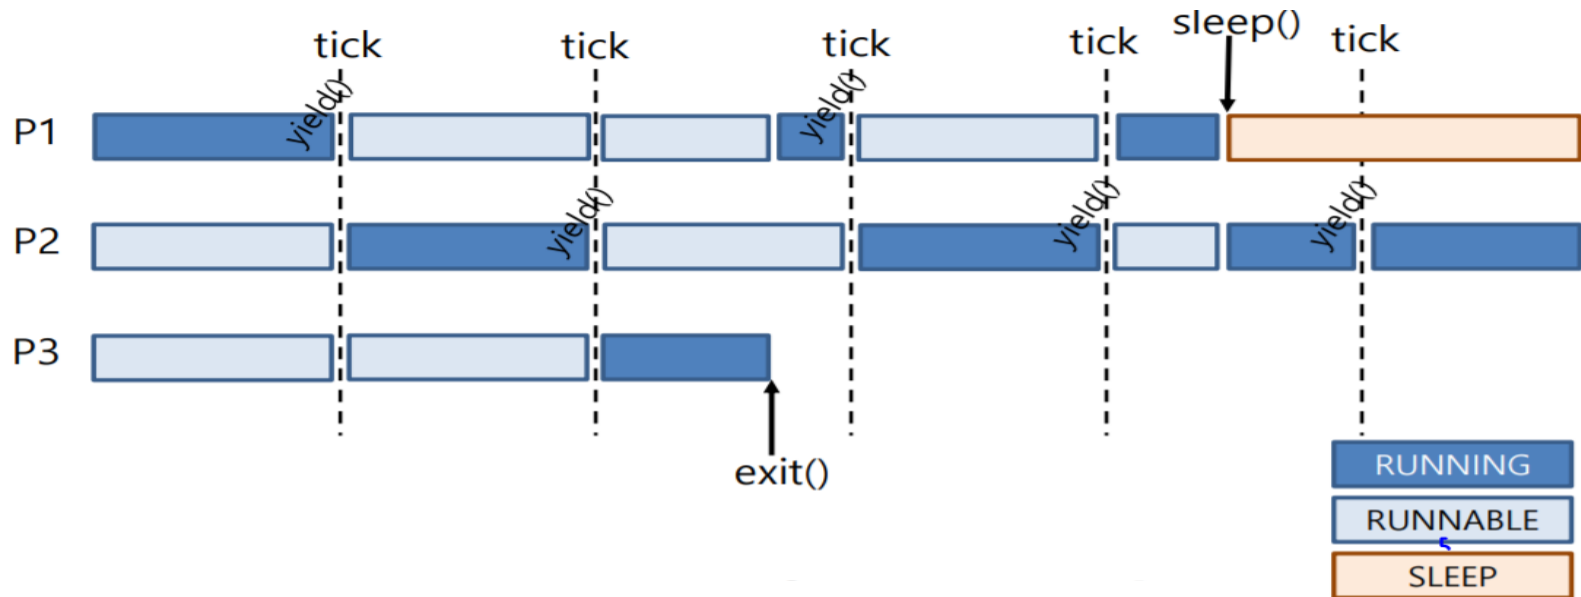
CPU scheduling

- Selects from the processes in memory that are ready to execute, and allocates CPU to one of them



How current scheduler works in xv6?

- Every timer IRQ enforces a yield of a CPU
- Process to be scheduled to be RUNNING state will be chosen in round-robin manner



Strawman scheduler

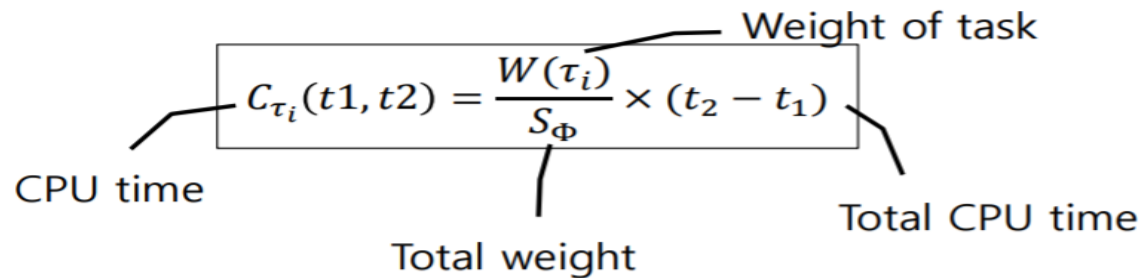
- Organize all processes as a simple list
- In `schedule()`:
 - Pick first one on a list to run next
 - Put suspended task at the end of the list
- Problems?
 - Allows only round-robin scheduling
 - Can't prioritize tasks

Fair scheduling

- And, how should time slices be distributed according to priority?
 - The difference of time slice by the nice value is not fair
 - E.g, processes with nice value 20 and 21 are given 100ms and 95ms
 - Processes with nice value 38 and 39 are given 10ms and 5ms
 - The differences are same to 5ms, but it's not proportional
 - To solve this problem, CFS(Completely Fair Scheduler) has been used since Linux kernel 2.6.23.

CFS (Completely Fair Scheduling)

- Linux default scheduler
- Basic concept
 - The CPU is allocated to the process in proportion to its weight
 - CPU time of any task satisfies in any given time between t_1 and t_2



The diagram shows the formula for CPU time allocation in CFS: $C_{\tau_i}(t_1, t_2) = \frac{W(\tau_i)}{S_\Phi} \times (t_2 - t_1)$. Hand-drawn lines point from text labels to parts of the formula: 'CPU time' points to the left side of the equation; 'Total weight' points to the denominator S_Φ ; 'Weight of task' points to the numerator $W(\tau_i)$; and 'Total CPU time' points to the term $(t_2 - t_1)$.

$$C_{\tau_i}(t_1, t_2) = \frac{W(\tau_i)}{S_\Phi} \times (t_2 - t_1)$$

- Nice to weight
 - Difference in nice by 1 provides 10% more (or less) CPU time
 - However, the larger the absolute value of nice, the smaller the ratio between the two values
 - Therefore, a new concept “weight”
 - Although there is formula, hard-code pre-defined array like Linux

$$weight = 1024(weight\ of\ nice\ 20) \times (1.25)^{-(nice-20)}$$

CFS parameters

- Time slice

- Task's minimum time to be executed before it is preempted
- Allocated to the process in proportion to its weight

$$time_slice = scheduling_latency \times \frac{weight\ of\ task}{total\ weight\ of\ runqueue}$$

- Scheduling latency (6ms by default)
 - Minimum time period to satisfy proportional CPU time distribution

- vruntime (virtual runtime)

- Accounts for how long a process has run proportional to its weight
- It's easy to compare how fairly the CPU is allocated
- By comparing this value, you can select the next process to be scheduled

$$vruntime = (actual\ runtime) \times \frac{weight\ of\ nice\ 20\ (1024)}{weight\ of\ task}$$

CFS scheduling

1. A task with minimum virtual runtime is scheduled
2. Scheduled task gets time slice proportional to its {weight / total weight}
3. While the task is running, virtual runtime is updated
4. After task runs more than time slice, go back to 1

Project 3. Implement CFS on xv6

- Implement CFS on xv6

- Select process with minimum virtual runtime from runnable processes
- Update runtime/vruntime for each timer interrupt
- If task runs more than time slice, enforce a yield of the CPU
- Default nice value is 20, ranging from 0 to 39, and weight of nice 20 is 1024
- Nice(0~39) to weight(Although there is formula, hard-code pre-defined array like Linux)

$$weight = \frac{1024}{(1.25)^{nice-20}}$$

/* 0 */	88761,	71755,	56483,	46273,	36291,
/* 5 */	29154,	23254,	18705,	14949,	11916,
/* 10 */	9548,	7620,	6100,	4904,	3906,
/* 15 */	3121,	2501,	1991,	1586,	1277,
/* 20 */	1024,	820,	655,	526,	423,
/* 25 */	335,	272,	215,	172,	137,
/* 30 */	110,	87,	70,	56,	45,
/* 35 */	36,	29,	23,	18,	15,

- Time slice calculation (our scheduling latency is 10ticks)

$$time\ slice = 10tick \times \frac{weight\ of\ current\ process}{total\ weight\ of\ runnable\ processes}$$

- vruntime calculation

$$vruntime += \Delta runtime \times \frac{weight\ of\ nice\ 20\ (1024)}{weight\ of\ current\ process}$$

Project 3. Implement CFS on xv6

- How about newly forked process?
 - A process inherits the parent process's vruntime
- How about woken process?
 - When a process is woken up, its virtual runtime gets
(minimum vruntime of processes in the ready queue – $vruntime(1tick)$)
$$vruntime(1tick) = 1tick \times \frac{\text{weight of nice } 20 (1024)}{\text{weight of current process}}$$

(If there is no process in the RUNNABLE state when a process wakes up, you can set the vruntime of the process to be woken up to “0”)
- DO NOT call sched() during a wake-up of a process
 - Ensure that the time slice of the current process expires
 - Woken-up process will have the minimum vruntime (by the formula above)
 - But we do NOT want to schedule the woken-up process before the time slice of current process expires
 - This is by default in xv6

Project 3. Implement CFS on xv6

- To check if CFS is implemented properly, `ps()` should be modified
- Expected output (`mytest.c`)

```
$ mytest
=== TEST START ===
name      pid      state      priority    runtime/weight  runtime      vruntime      tick 4144000
init       1        SLEEPING    5           0              2000        1000
sh         2        SLEEPING    5           0              0           0
mytest     3        RUNNABLE    5          25             752000      35935
mytest     4        RUNNING     0          35            3122000     35903
```

- Print out the following information about the processes
- **Use millitick unit** (multiply the tick by 1000)
 - **runtime, vruntime, total tick**
 - Do NOT use float/double types to present runtime and vruntime
 - Kernel avoid floating point operation as much as possible
- Indents of name section should be aligned even if process has long name (up to 10 letters) or very large value... (runtime, vruntime)

Project 3. Implementation details

- Project 3 should be done based on your project 2 code
- Never use float or double types to represent runtime and vruntime.
- Consider the case of integer overflow vruntime
 - Even if over the scope of integer, shall operate without problems
 - And it must be printed normally
 - Do not worry about runtime, total tick
- FAQ
 - Q : My time slice is 6.5. However, what if timer interrupt occurs every 1 tick? (context switch can occur only with 1 tick)
 - A :Tasks will run over it's time slice (7 ticks) & add vruntime

Submission

- Please implement CFS on xv6
- Use the ***submit & check-submission*** binary file in Ji Server
 - make clean
 - \$ ~swe3004/bin/submit pa3 xv6-public
 - you can submit several times, and the submission history can be checked through check-submission
 - Only the last submission will be graded

Submission

- PLEASE DO NOT COPY
 - We will run inspection program on all the submissions
 - Any unannounced penalty can be given to **both students**
 - 0 points / negative points / F grade ...
- Due date: 10/18(Wed.), 23:59:59 PM
 - -25% per day for delayed submission

Questions

- If you have questions, please ask on i-campus discussion section
 - Please use the discussion board
 - Discussion board preferred over messages
- You can also visit Corporate Collaboration Center #85533
 - Please iCampus message TA before visiting
- Reading xv6 commentary will help you a lot
 - <http://csl.skku.edu/uploads/SSE3044S20/book-rev11.pdf>