

SWE3004 Operating Systems, fall 2023

Project 6. File extension

TA)

Jongseok Kim

Shinhyun Park

Hyeonmyeong Lee

Gwanjong Park

Project plan

- Total 6 projects

- ~~1) Booting xv6 operating system~~

- ~~2) System call~~

- ~~3) CPU scheduling~~

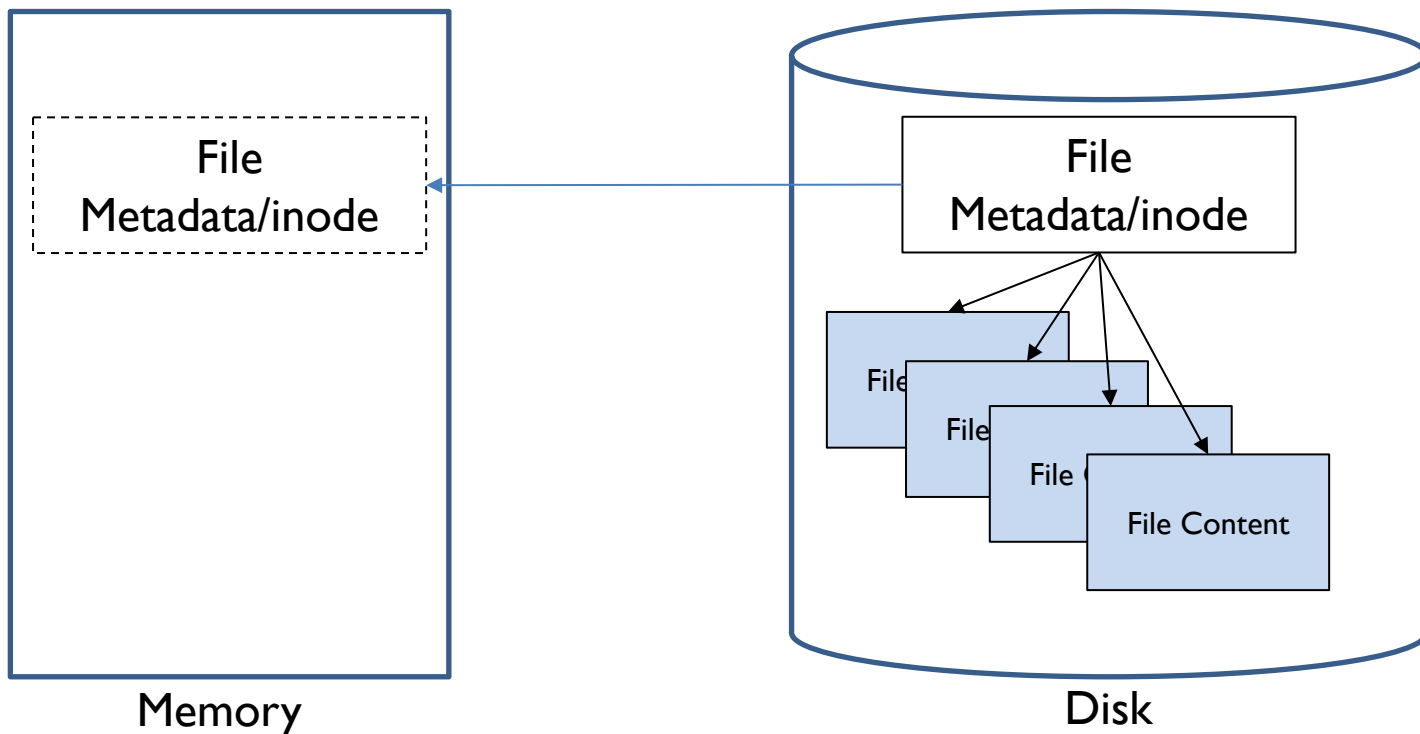
- ~~4) Virtual memory~~

- ~~5) Page replacement~~

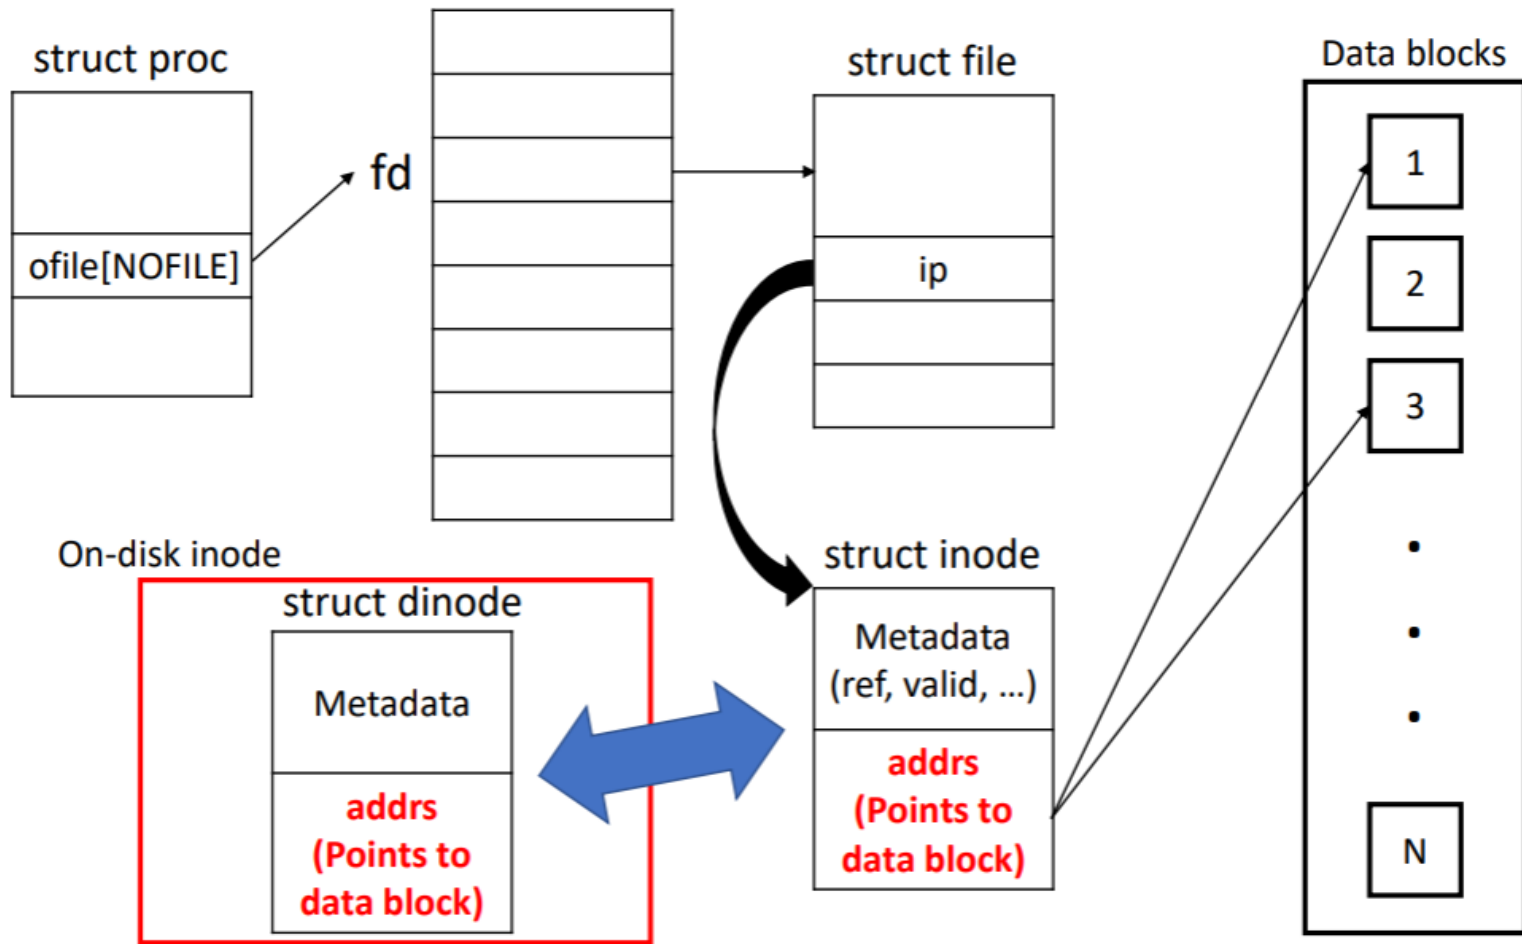
- 6) File systems (Opt, optional)

How open file works

- For each file, we have
 - File contents (data)
 - File attributes (metadata or inode)




xv6 filesystem



xv6 filesystem: superblock

- Disk layout of xv6 filesystem



- Contents of struct superblock are decided at “mkfs”

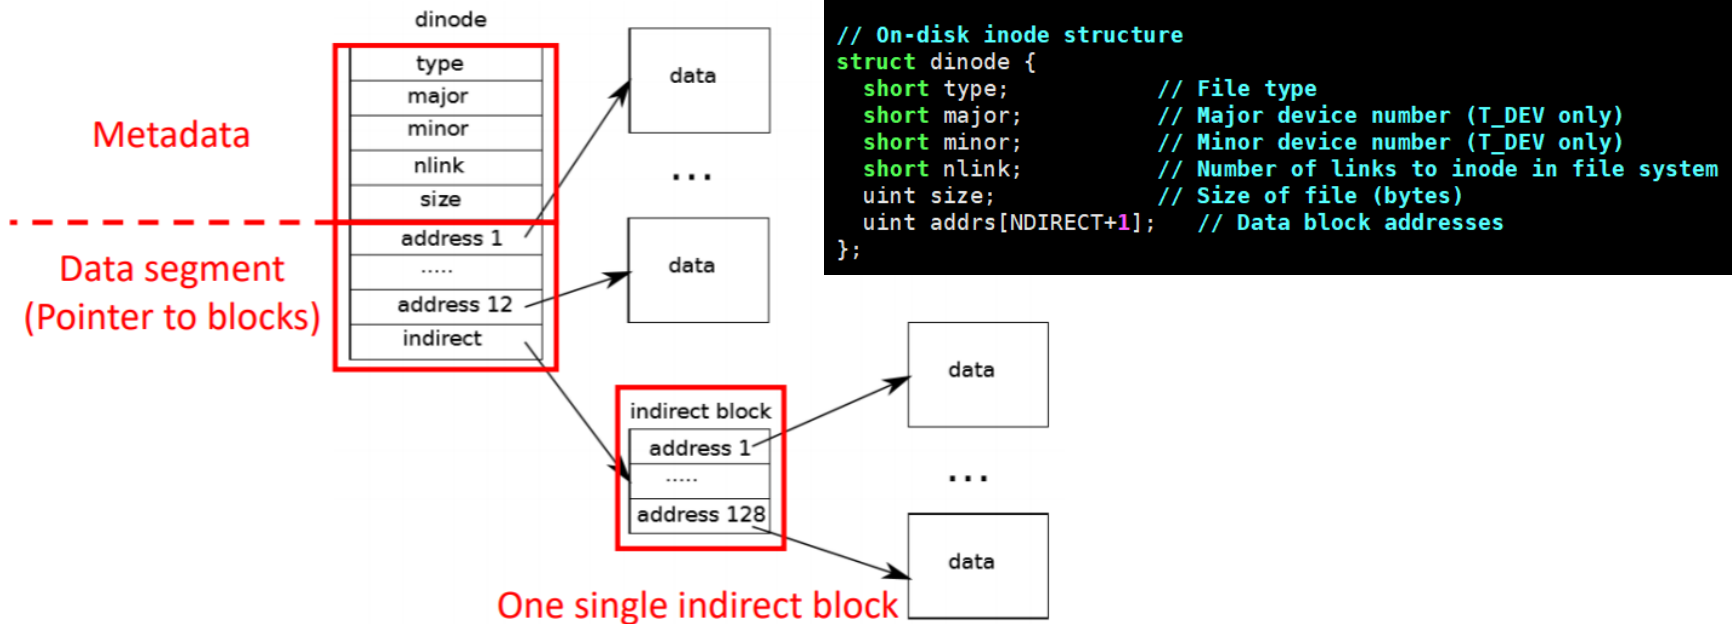
```
struct superblock {  
    uint size;           // Size of file system image (blocks)  
    uint nblocks;        // Number of data blocks  
    uint ninodes;        // Number of inodes.  
    uint nlog;           // Number of log blocks  
    uint logstart;       // Block number of first log block  
    uint inodestart;     // Block number of first inode block  
    uint bmapstart;      // Block number of first free map block  
};
```

<fs.h>

```
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap start 58
```

xv6 filesystem: inodes

- On-disk inode is defined by *struct dinode* in fs.h



<Representation of file on disk of xv6>

xv6 filesystem: bmap() function

- The function bmap() manages the representation
 - bmap() is called both when reading and writing a file
 - When writing, bmap() allocates new blocks as needed to hold file content, as well as allocating an indirect block if needed to hold block addresses

```
static uint
bmap(struct inode *ip, uint bn)
{
    uint addr, *a;
    struct buf *bp;

    if(bn < NDIRECT){
        if((addr = ip->addrs[bn]) == 0)
            ip->addrs[bn] = addr = balloc(ip->dev);
        return addr;
    }
    bn -= NDIRECT;

    if(bn < NINDIRECT){
        // Load indirect block, allocating if necessary.
        if((addr = ip->addrs[NDIRECT]) == 0)
            ip->addrs[NDIRECT] = addr = balloc(ip->dev);
        bp = bread(ip->dev, addr);
        a = (uint*)bp->data;
        if((addr = a[bn]) == 0){
            a[bn] = addr = balloc(ip->dev);
            log_write(bp);
        }
        brelse(bp);
        return addr;
    }

    panic("bmap: out of range");
}
```

NDIRECT blocks are listed in inode itself

On-demand allocation for not allocated blocks

Project 6


- In this project, you have to implement following two features in xv6 file system
 - A. File extension
 - B. Block group
- You should implement both features on one initial xv6 source code

A. File extension in xv6

- Increase the maximum size of an xv6 file
 - Currently xv6 files are limited to 140 sectors (70 kilobytes)
 - 12 “direct” blocks and one “singly-indirect” block
 - Modify the system to support “doubly-indirect” block
 - 11 “direct” blocks, one “singly-indirect” block and one “doubly-indirect block)
 - The result will be that a file will be able to consist of up to 16523 sectors (about 8.5 megabytes)
- Codes to be modified
 - Some header files (fs.h, file.h... as you need)
 - bmap() function (in fs.c)
 - mkfs parameters (in mkfs.c / param.h)

```
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap start 58
```

```
sb: size 21113 nblocks 21049 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap start 58
```



- Use template codes given in Project 0,1,2 (with 1 CPU)
- Test code to check maximum file size will be given


wrote 140 sectors → wrote 16523 sectors

B. Block Group in xv6

- Block-group is literally a group of blocks that consists of inode blocks, bitmap blocks, data blocks
- With the introduction block group, data and inode blocks are evenly distributed across the physical disk
- Without block group, file system will be look like:



- But with block group, file system will be look like:



B. Block Group in xv6 (cnt'd)

- You should determine a block group size (BGSIZE) in accordance with FSSIZE
 - By default, BGSIZE equals to $FSSIZE/32$
 - But minimum BGSIZE should be at least 4096 (blocks)
 - If FSSIZE is not divisible by 32, remaining blocks should be added to log area
 - # of inode blocks should be $BGSIZE/4$
 - If BGSIZE is not divisible by 4, remaining blocks should be added to data blocks
 - You should determine the # of bitmap blocks so that all data blocks can be represented within bitmap blocks
 - Remaining blocks in a block group are all data blocks